

```
In [95]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [97]: df = pd.read_csv("emails.csv")


# Rename column if needed
if 'spam_or_not' in df.columns:
    df.rename(columns={'spam_or_not': 'spam'}, inplace=True)

df.head()
```

```
Out[97]:
```

	text	spam
0	Subject: naturally irresistible your corporate...	1
1	Subject: the stock trading gunslinger fanny i...	1
2	Subject: unbelievable new homes made easy im ...	1
3	Subject: 4 color printing special request add...	1
4	Subject: do not have money , get software cds ...	1

```
In [99]: plt.figure(figsize=(5,4))
sns.countplot(x='spam', data=df, hue='spam', palette=['green', 'red'], legend=False)
plt.xticks([0,1], ['Not Spam', 'Spam'])
plt.title("Spam vs Not Spam Distribution")
plt.xlabel("Class")
plt.ylabel("Count")
plt.show()
```



```
In [101... X = df['text']
y = df['spam']
```

```
In [103... X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.3, random_state=0  
)
```

```
In [105... vectorizer = TfidfVectorizer(stop_words='english')  
  
X_train_vec = vectorizer.fit_transform(X_train)  
X_test_vec = vectorizer.transform(X_test)
```


```
In [106... model = MultinomialNB()  
model.fit(X_train_vec, y_train)  
  
y_pred = model.predict(X_test_vec)
```

```
In [109... cm = confusion_matrix(y_test, y_pred)  
print("Confusion Matrix:\n", cm)
```

Confusion Matrix:

```
[[1314  0]  
[ 159 246]]
```

```
In [111... plt.figure(figsize=(5,4))  
sns.heatmap(  
    cm,  
    annot=True,  
    fmt='d',  
    cmap='Greens',  
    xticklabels=['Not Spam', 'Spam'],  
    yticklabels=['Not Spam', 'Spam'])  
  
plt.title("Confusion Matrix - Spam Detection")  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.show()
```



```
In [113... TN, FP, FN, TP = cm.ravel()  
  
print("True Positive (TP) - Spam correctly predicted:", TP)  
print("False Positive (FP) - Not Spam predicted as Spam:", FP)
```

```
print("True Negative (TN) - Not Spam correctly predicted:", TN)
print("False Negative (FN) - Spam predicted as Not Spam:", FN)
```

```
True Positive (TP) - Spam correctly predicted: 246
False Positive (FP) - Not Spam predicted as Spam: 0
True Negative (TN) - Not Spam correctly predicted: 1314
False Negative (FN) - Spam predicted as Not Spam: 159
```

```
In [115]: accuracy = accuracy_score(y_test, y_pred)
error_rate = 1 - accuracy

print("\nAccuracy:", accuracy)
print("Error Rate:", error_rate)
```

```
Accuracy: 0.9075043630017452
Error Rate: 0.0924956369982548
```

```
In [117]: print("\nClassification Report:\n")
print(classification_report(
    y_test,
    y_pred,
    target_names=['Not Spam', 'Spam']
))
```

```
Classification Report:
```

	precision	recall	f1-score	support
Not Spam	0.89	1.00	0.94	1314
Spam	1.00	0.61	0.76	405
accuracy			0.91	1719
macro avg	0.95	0.80	0.85	1719
weighted avg	0.92	0.91	0.90	1719

```
In [ ]:
```